sample_01 vs sample_02

Generated 2026-08-22 00:09 · 700 images

What should I look at?

186 images have low semantic fit. They don't strongly match any of the detected themes — worth a quick look to confirm they belong in this dataset.

26 images belong to near-duplicate groups. This may be expected (e.g. burst-mode photos, or the same subject appearing in both datasets) or a sign of redundancy, depending on what this dataset is for.

Compared with “sample_02”: distributional distance 0.152. For reference, splitting this dataset into two random halves gives 0.032 — that's roughly the gap you'd expect from sampling alone, even between two halves of the exact same dataset. A distance well above that suggests a genuinely different overall makeup; a distance close to it suggests the two are hard to tell apart, distribution-wise.

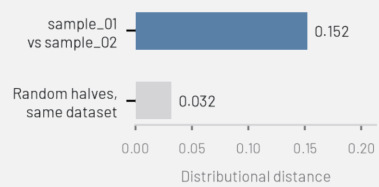

Texture distance from “sample_02”: 0.083. This looks at low-level visual texture (not subject matter) — splitting this dataset into two random halves gives 0.077 as a reference. Two datasets can be texturally similar even when the distributional distance above says they differ in subject matter, or vice versa — different cameras, compression, or rendering can show up here even when the content itself looks the same.

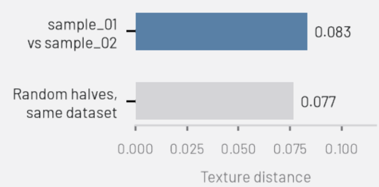

OVERVIEW

700
TOTAL IMAGES
500 200

73%
SEMANTIC COVERAGE
81% 54%

26.6%
LOW SEMANTIC FIT
18.8% 46.0%

3.3%
VISUAL REDUNDANCY
4.2% 1.0%

10
SEMANTIC AXES

14% / 6%
LARGEST / SMALLEST CLUSTER

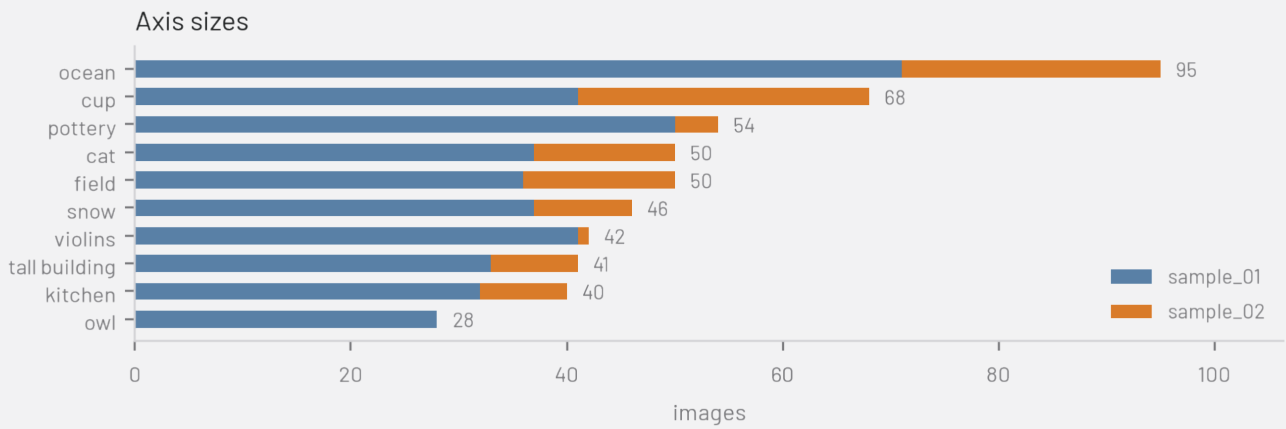

Axis sizes
ocean
cup
pottery
cat
field
snow
violins
tall building
kitchen
owl
95
68
54
50
50
46
42
41
40
28
sample_01
sample_02
0
20
40
60
80
100
images

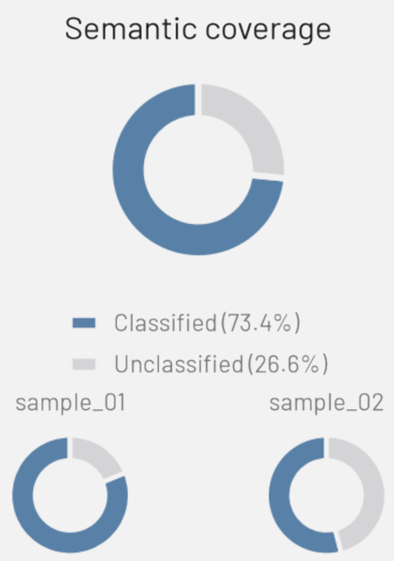

Semantic coverage
Classified (73.4%)
Unclassified (26.6%)
sample_01
sample_02

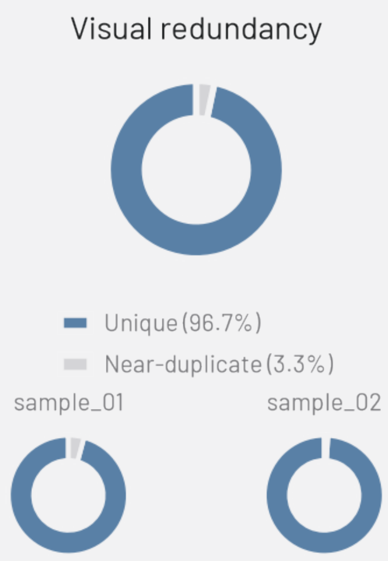

Visual redundancy
Unique (96.7%)
Near-duplicate (3.3%)
sample_01
sample_02

WHAT IS IN YOUR DATASETS?

One representative image per semantic axis — for auto-detected axes, the image closest to the cluster centroid; for custom text axes (which have no centroid image of their own), a random image that dominates that axis.

field

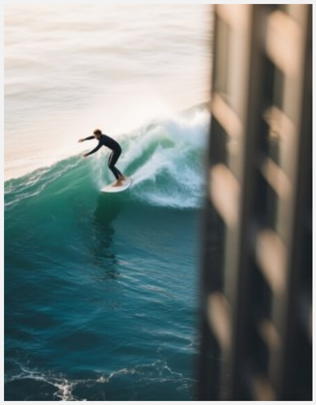

ocean

cup

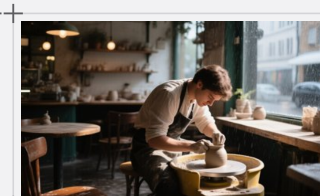

pottery

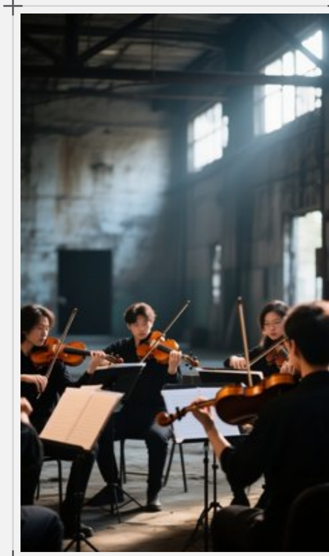

violins

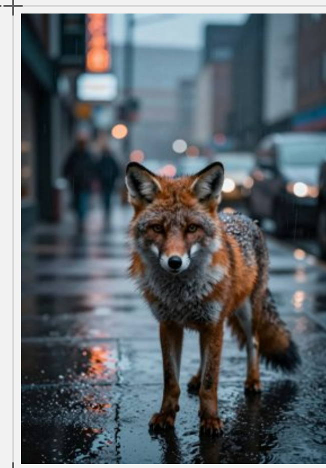

snow

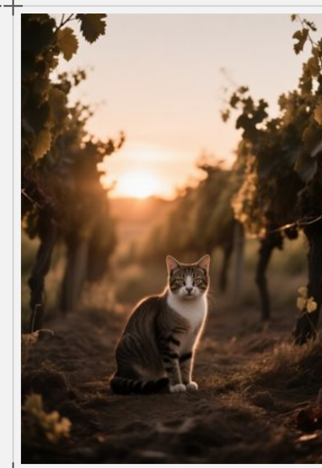

cat

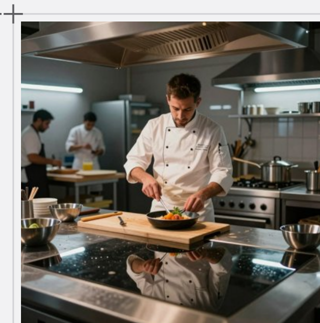

kitchen

owl

tall building

DATASET COMPOSITION

Two independent signals: whether an image fits a semantic theme at all, and whether it's visually a near-duplicate of another image.

What doesn't fit?: Low semantic fit — 186 images (26.6%)

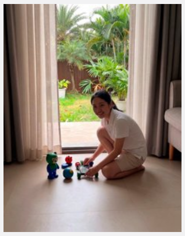

0692.jpg
sample_01

0334.jpg
sample_01

0125.jpg
sample_01

0280.jpg
sample_01

0230.jpg
sample_01

+181
MORE...

How much visual redundancy is there?: Near duplicates — 26 images (3.7%)

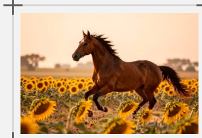

Group 1
0680.jpg
sample_01

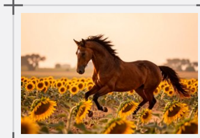

0681.jpg
sample_01

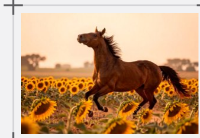

0683.jpg
sample_01

0065.jpg
sample_02

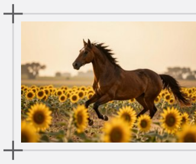

0165.jpg
sample_02

+2
MORE...

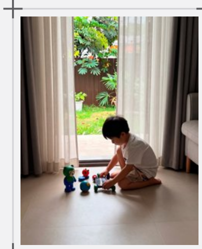

Group 2
0687.jpg
sample_01

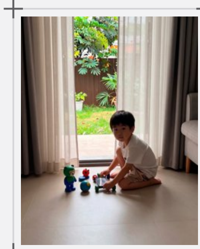

0688.jpg
sample_01

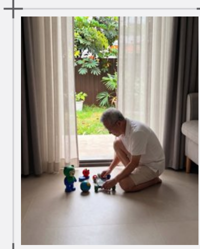

0689.jpg
sample_01

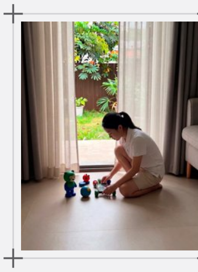

0690.jpg
sample_01

0096.jpg
sample_02

+2
MORE...

Group 3
0122.jpg
sample_01

0685.jpg
sample_01

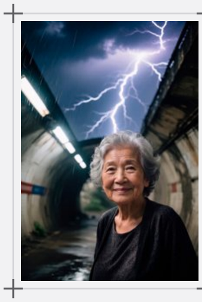

0686.jpg
sample_01

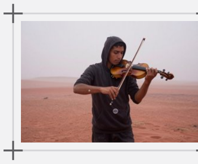

Group 4
0657.jpg
sample_01

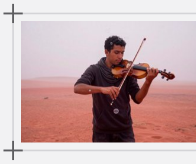

0695.jpg
sample_01

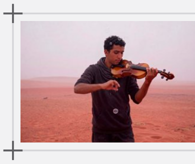

0696.jpg
sample_01

Group 5
0559.jpg
sample_01

0699.jpg
sample_01

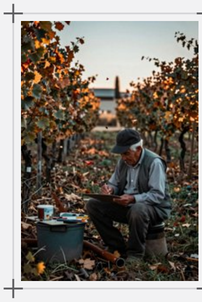

Group 6
0618.jpg
sample_01

0697.jpg
sample_01

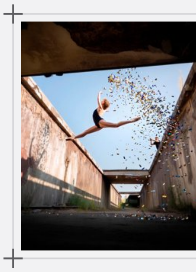

Group 7
0693.jpg
sample_01

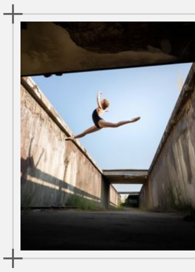

0051.jpg
sample_02

HOW DIVERSE ARE THEY?

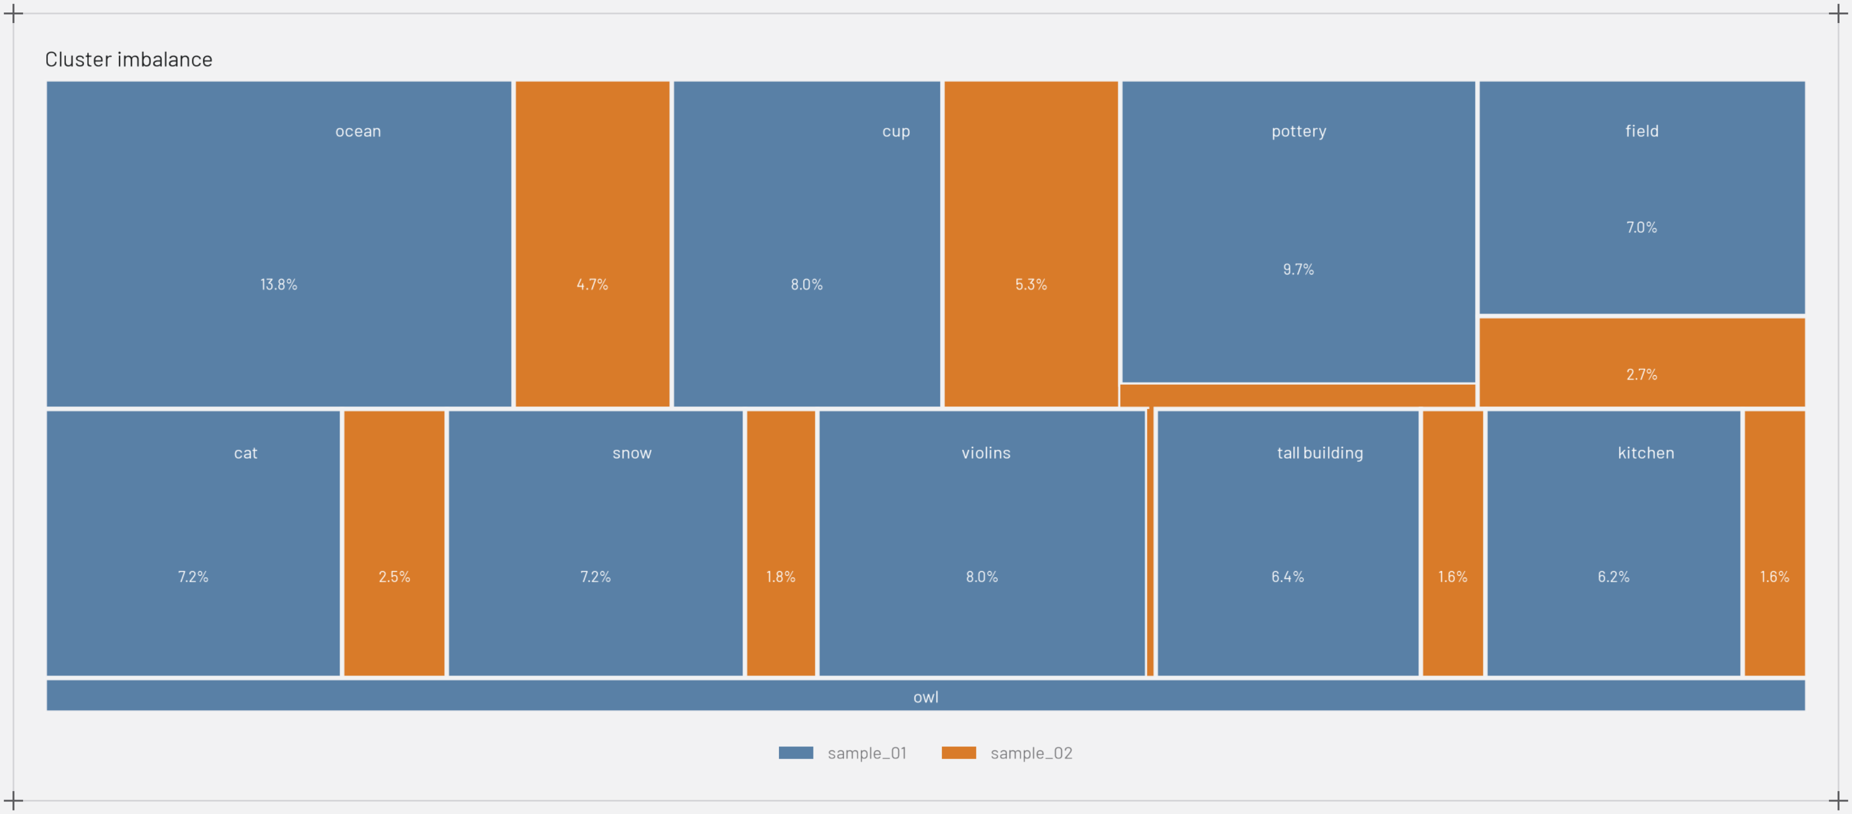

HOW DO THE TWO DATASETS COMPARE?

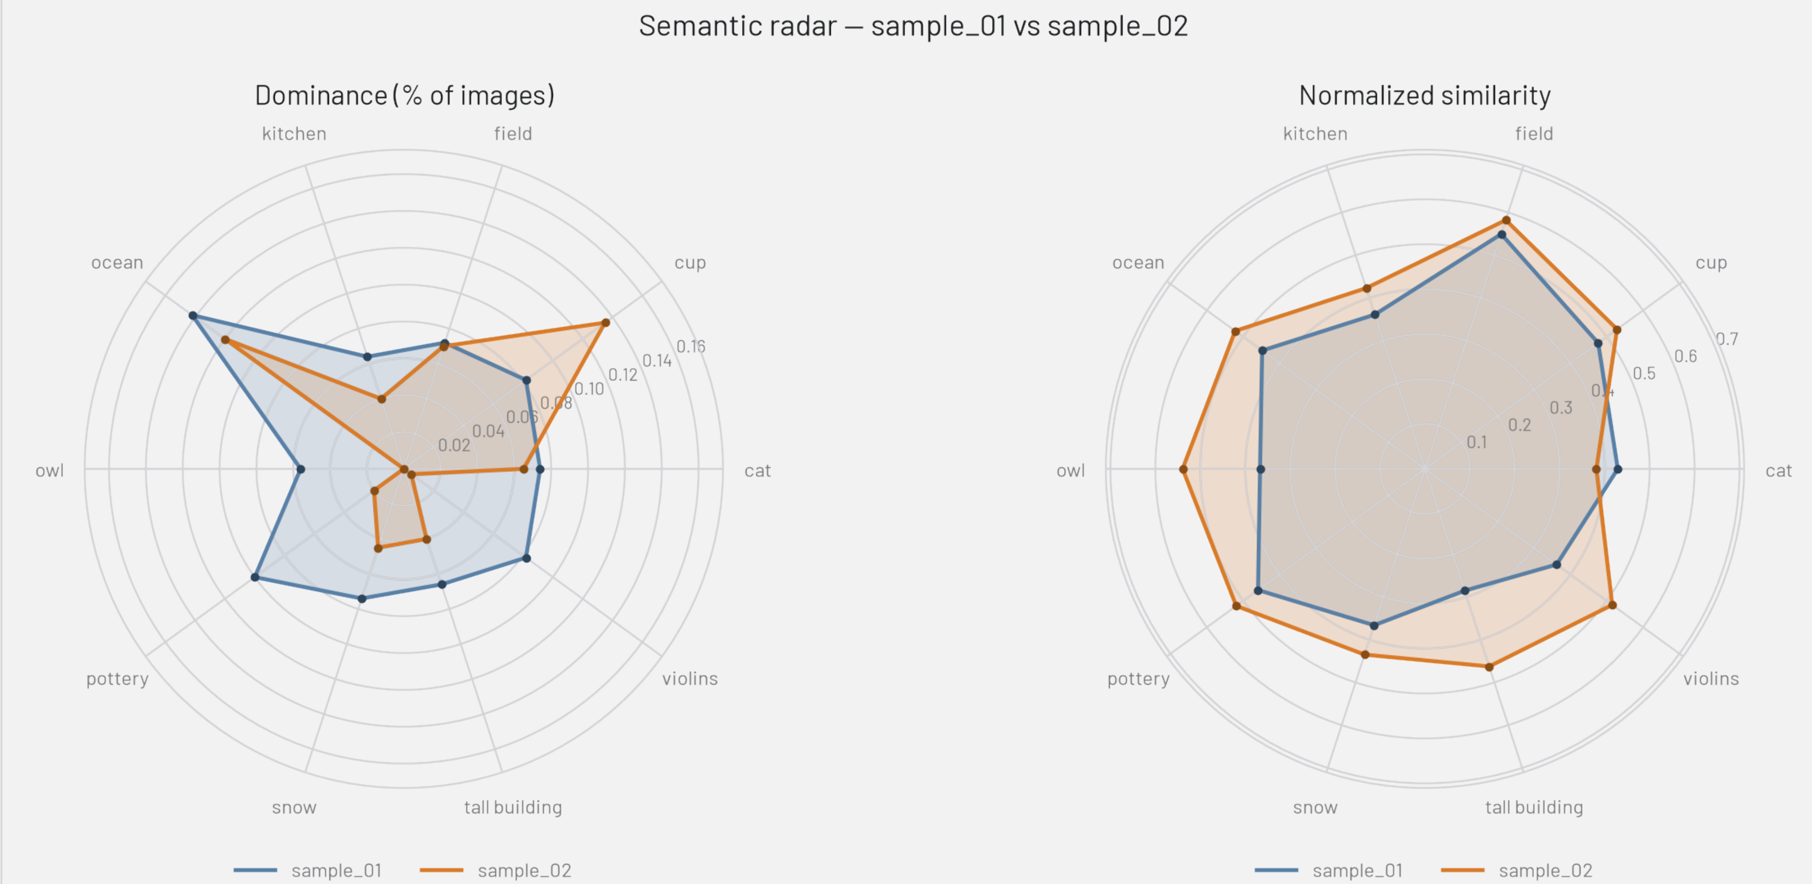


UMAP — sample_01 vs sample_02
Axis
Other
cat
cup
field
kitchen
ocean
owl
pottery
snow
tall building
violins
Dataset
sample_01
sample_02