sample_01

Generated 2026-08-20 16:32 · 480 images

What should I look at?

88 images have low semantic fit. They don't strongly match any of the detected themes — worth a quick look to confirm they belong in this dataset.

Compared with “sample_02”: distributional distance 0.155. For reference, splitting this dataset into two random halves gives 0.000 — that's roughly the gap you'd expect from sampling alone, even between two halves of the exact same dataset. A distance well above that suggests a genuinely different overall makeup; a distance close to it suggests the two are hard to tell apart, distribution-wise.

OVERVIEW

480	82%	18.3%
TOTAL IMAGES	SEMANTIC COVERAGE	LOW SEMANTIC FIT

0.0%	10	14% / 6%
VISUAL REDUNDANCY	SEMANTIC AXES	LARGEST / SMALLEST CLUSTER

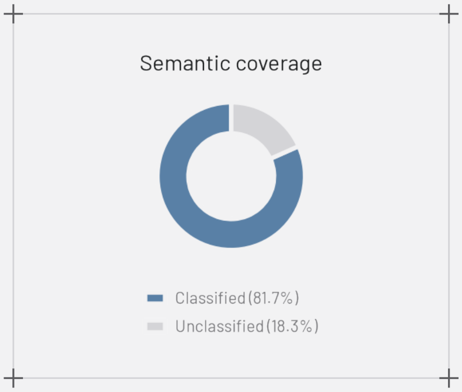

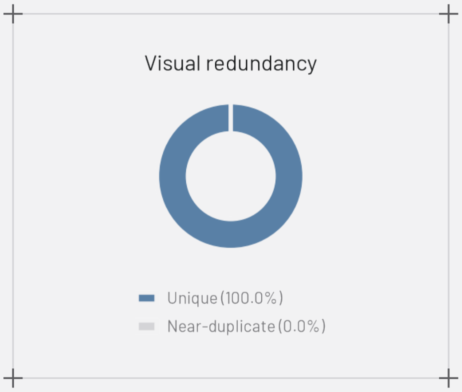

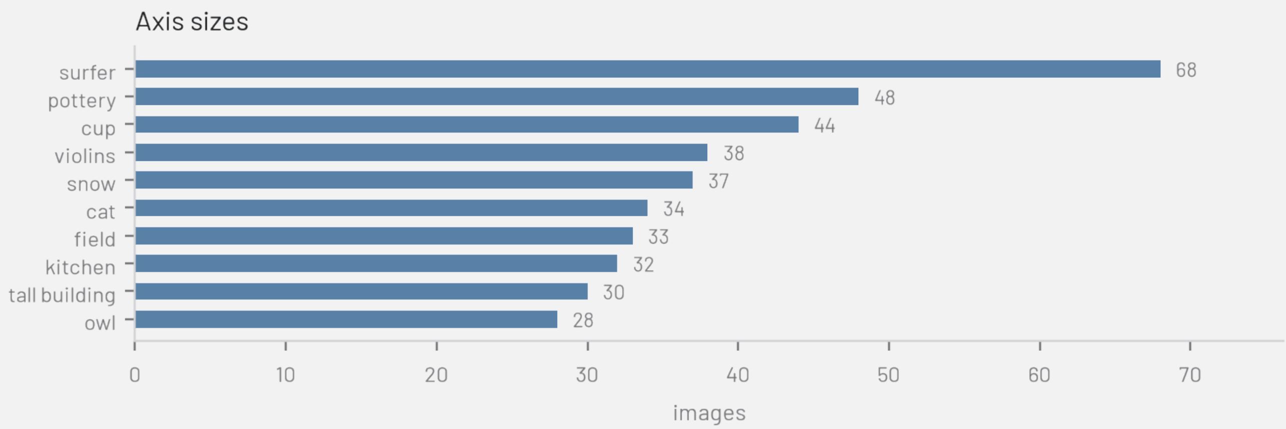

WHAT'S IN THIS DATASET

One representative image per semantic axis — for auto-detected axes, the image closest to the cluster centroid; for custom text axes (which have no centroid image of their own), a random image that dominates that axis.

field

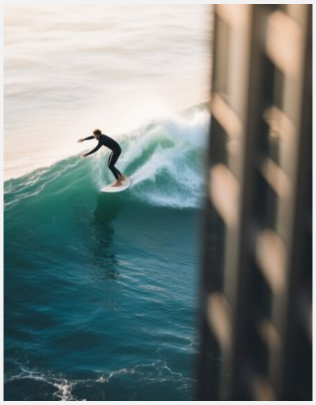

surfer

cup

pottery

snow

violins

cat

kitchen

owl

tall building

DATASET COMPOSITION

Two independent signals: whether an image fits a semantic theme at all, and whether it's visually a near-duplicate of another image.

Low semantic fit — 88 images (18.3%)

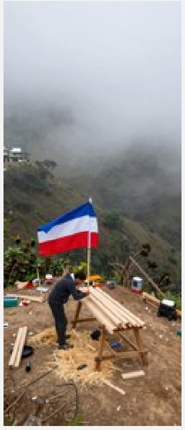

0334.jpg

0125.jpg

0636.jpg

0280.jpg

0230.jpg

0058.jpg

Near duplicates — 0 images (0.0%)

How diverse is it? — Cluster imbalance

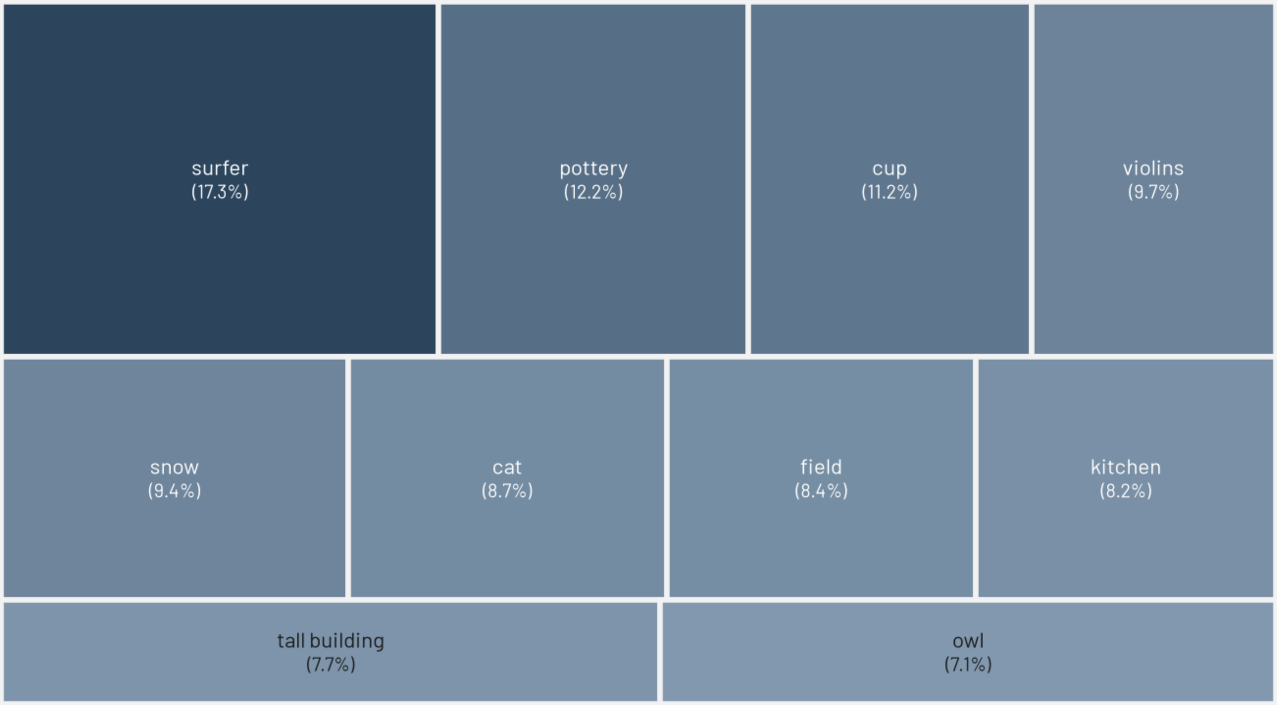

Semantic radar — sample_01

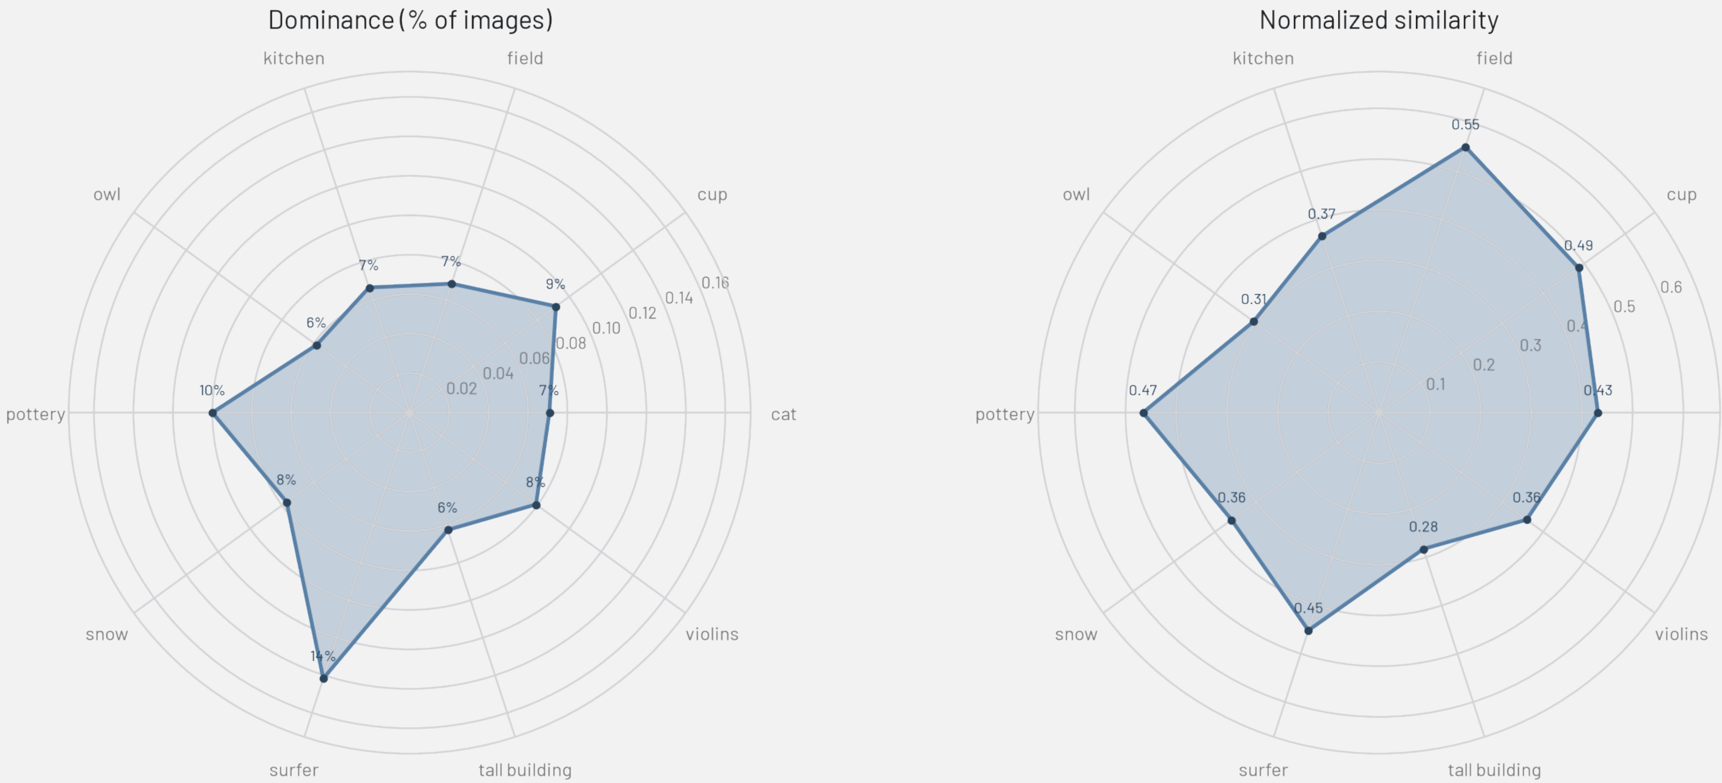

UMAP — sample_01