

sample_01 vs sample_02

Generated 2026-08-22 16:25 · 700 images



What should I look at?

38 images have low semantic fit. They don't strongly match any of the detected themes — worth a quick look to confirm they belong in this dataset.

26 images belong to near-duplicate groups. This may be expected (e.g. burst-mode photos, or the same subject appearing in both datasets) or a sign of redundancy, depending on what this dataset is for.

Compared with “sample_02”: distributional distance 0.152. For reference, splitting this dataset into two random halves gives 0.022 — that’s roughly the gap you’d expect from sampling alone, even between two halves of the exact same dataset. A distance well above that suggests a genuinely different overall makeup; a distance close to it suggests the two are hard to tell apart, distribution-wise.

Texture distance from “sample_02”: 0.083. This looks at low-level visual texture (not subject matter) — splitting this dataset into two random halves gives 0.000 as a reference. Two datasets can be texturally similar even when the distributional distance above says they differ in subject matter, or vice versa — different cameras, compression, or rendering can show up here even when the content itself looks the same.



OVERVIEW

700

TOTAL IMAGES

500

200

95%

SEMANTIC COVERAGE

92%

100%

5.4%

LOW SEMANTIC FIT

7.6%

0.0%

3.3%

VISUAL REDUNDANCY

4.2%


1.0%

20


SEMANTIC AXES

13% / 0%


LARGEST / SMALLEST CLUSTER



Semantic coverage



Visual redundancy



WHAT IS IN YOUR DATASETS?


One representative image per semantic axis – for auto-detected axes, the image closest to the cluster centroid; for custom text axes (which have no centroid image of their own), a random image that dominates that axis.



open fire



water



waves



clothes



cup




cat



violins




snow




chef



horse



owl



tall building



astronaut



lavender fields



road



pottery




train



air



tunnel




toys


DATASET COMPOSITION

Two independent signals: whether an image fits a semantic theme at all, and whether it's visually a near-duplicate of another image.

What doesn't fit?: Low semantic fit — 38 images (5.4%)



How much visual redundancy is there?: Near duplicates — 26 images (3.7%)





Group 3
0122.jpg
sample_01



0685.jpg
sample_01



0686.jpg
sample_01




Group 4
0657.jpg
sample_01




0695.jpg
sample_01



0696.jpg
sample_01



Group 5
0559.jpg
sample_01




0699.jpg
sample_01




Group 6
0618.jpg
sample_01



0697.jpg
sample_01




Group 7
0693.jpg
sample_01





0051.jpg
sample_02

HOW DIVERSE ARE THEY?



HOW DO THE TWO DATASETS COMPARE?

Semantic radar — sample_01 vs sample_02



UMAP – sample_01 vs sample_02

